

# Assignment #5: Synchronization

CS201 Spring 2020

10 points

due Sunday, Mar. 29th, 11:59 pm

second-chance (for 9 points) due Sunday, Apr. 5th, 11:59 pm

## 1 Synchronization

You'll write a multithreaded C program that will solve the Readers/Writers Problem of the First Kind, as we discussed in class.

### 1.1 Algorithm

Here's the complete solution, in pseudocode:

```
// initialization
semaphore rw_sem = 1;
semaphore mutex = 1;
int num_readers = 0;

// reader thread does this
for (i=0; i<NUM_READS; ++i) {
    wait(mutex); // lock for accessing num_readers
    ++num_readers;
    // print a message with the current value of num_readers
    if (num_readers == 1) // if I am first reader
        wait(rw_sem); // lock for shared buffer
    signal(mutex); // done accessing num_readers
    // print a message that this reader is reading
    // perform reading
    wait(mutex); // lock for accessing num_readers
    --num_readers;
    // print a message with the current value of num_readers
    if (num_readers == 0) // if I am last reader
        signal(rw_sem); // release lock for shbuf
    signal(mutex); // done accessing num_readers
}

// writer thread does this
for (i=0; i<NUM_WRITES; ++i) {
    wait(rw_sem); // wait
    // print a message that this writer is writing
    // do the writing
    signal(rw_sem); // done
}
```

## 1.2 Named semaphores

On Linux and macOS, a named semaphore is a system-wide resource. The name lets a user uniquely identify a semaphore (so that my semaphores don't get mixed up with your semaphores).

The call `sem_open()` initializes a named semaphore. By specifying `O_CREAT` and `O_EXCL`, you can create a named semaphore so that only you can use it.

The call `sem_close()` tells the system that you're done using the semaphore, and `sem_unlink()` actually removes the semaphore.

If your program exits (or crashes) after `sem_open()` but before `sem_unlink()`, you'll get an error the next time you try to `sem_open()` it. In this case, clean up the semaphore with `sem_unlink()` and try again.

The example program `semaphore-example.c`, under Examples in the class gitlab space, shows how to create and use and clean up a named semaphore.

## 1.3 Declarations

You will have to do this assignment on Linux or macOS. You should run your code on kaladin though to verify that it works correctly.

Declare these structures:

```
typedef struct {
    pthread_mutex_t mutex;
    sem_t *rw_sem;
    int readerCount;
} SyncInfo;
```

```
typedef struct {
    SyncInfo *syncInfo;
    char myName[32];
} ThreadInfo;
```

You should write these two functions:

```
void *reader(void *data);
void *writer(void *data);
```

The `data` parameter will point to a `ThreadInfo` struct. Instead of the `while (true)`, the `reader()` function will have a for-loop in which it tries to perform reading. Similarly the `writer()` function will have a for-loop in which it tries to perform writing. The reading and writing actions will actually just be this call: `sleep(1)`.

## 1.4 Compiling and linking

You'll have to include these header files:

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <unistd.h>
#include <pthread.h>
#include <semaphore.h>
```

On Linux, you will have to compile your program this way:

```
$ gcc rw1.netid.c -lpthread
```

On macOS, you do not need the `-lpthread`.

## 2 What to Do

Do the actual implementations in C of the program described above. Put your structure definitions and function prototypes in `rw1.netid.h`, and put your code in `rw1.netid.c`; submit these two files to Blackboard.

You will need to write the `main()` function. Your `main()` will create the threads and then wait for them to finish. Leave your `main()` in your file when you submit it.

Create ten reader threads and two writer threads. Use these `#define` statements:

```
#define NUM_READERS 10
#define NUM_WRITERS 2
#define NUM_READS 5
#define NUM_WRITES 5
```

Define this also:

```
#define SEM_NAME "/SEM_netid"
```

and use your actual `netid`.

For an example of how to use a semaphore, see the sample program `semaphore-example.c`, in the Examples directory of the class gitlab site (<https://gitlab.uvm.edu/Jason.Hibbeler/ForStudents/tree/master/CS201/>).

Each thread will have its own `ThreadInfo` instance. Create an array of these, as you did in Assignment #4.

`NUM_READS` is the number of times that each reader goes through its loop, and `NUM_WRITES` is the number of times that each writer goes through its loop.

You need only a single instance of `SyncInfo`.

Give each thread a name. In the for-loop that you will use for creating the reader threads, do this:

```
sprintf(threadInfo[idx].myName, "Reader_%d", i);
```

Print out `hello from Reader n` for each reader when a reader thread starts; print a similar message for each writer rthread

Put your code in a file `rw1.netid.c` and your declarations in `rw1.netid.h`.

## 3 Testing

How will you know this is working correctly? You should see that only one writer at a time can write, but that more than one reader can read. Print a message each time a reader thread modifies the `numReaders` variable. This output will help you see the program is correct.